

Title

Davide Scassola, Salvatore Milite

January 19, 2019

The aim of the project is to implement a Binary Tree in c++ and compare its performance with `std::map`. The work is organized in three directories: in the *include* directory there is the templated class, in the *test* one there is the unit test and the generated documentation is in *documentation*. There is also an *old* with an early iterative version of the code. The benchmarks where we compare our version of the data structure against `std::map` and where we test our balance algorithm is in the main folder.

Code and implementation

The `BinaryTree` class is templated with three arguments, the key type, the value type and the function used for comparison, in particular the default is `std::less<value_type(key)>`, a functional object that returns, when it is called with its two arguments, *arg1* < *arg2*. The class has a `unique_pointer` to the root node and an instance of the comparison callable object. Throughout the code we decided to use smart pointers to make easier the memory management and to avoid to write a custom destructor.

Node class

The node data structure is nested into the class, we opted for this choice in order to preserve a coherence between the templates as well as a logical dependency (indeed we do not have to instantiate node objects outside of the class). The Node object is constructed with those attributes

- a `unique_pointer` to the right node
- a `unique_pointer` to the left node
- a pointer to the parent (this pointer does not have to release memory)
- a `std::pair` with const key and value

Iterators

The class is equipped with two types of iterators: const and non const, defined as internal classes with a pointer to a node. The iterators have overloaded the operators: `*`, which returns a `std::pair<>` from the node; `++`, which returns an iterator to the next node (next with respect to the order of the keys) and `++(int)` with the usual behavior.

Methods

The `BinaryTree` class has implemented the copy behavior (it creates a deep copy of tree with a recursive algorithm) and the move semantic; it also has a default constructor which initializes a void tree or can take the functional object as an argument (the destructor is the default). Regarding the other method we tried to make code as simple as possible, so we did a wide use of recursion, and most of the algorithms are recursive. The functions

`insert()` and `find()` use a common private method named `search()` which returns, given a key and a pair of node pointers, a reference to the node that matches the key or to the node corresponding to the insert position of that key. Actually, `search` returns a `std::pair` with a reference to a node pointer (the one described before) and a pointer, the latter is the pointer to the parent. This choice has been made as to remove some if statements (the root can be treated as an ordinal node) and make easier to assign the parent node. Indeed, when we assign the parent, if the node is on the right, we consider as its parent its parent's parent (i.e. the first greater parent). In these way the `++()` operator is easier to implement. The `insert()` method can work in two scenarios:

- the key is not present; in this case the function inserts the node and returns an `std::pair` with an iterator to the new node and a `bool` with the value `true`
- the key is present; in this case the function returns an `std::pair` consisting of an iterator to the node with the same key used for the search and a `bool` with the value `false`. Is up to the user the final choice of changing the value inside.

The `balance()` function is recursive. It first transforms the tree into a vector of `std::pairs` and after it constructs the balanced tree splitting the list and inserting the node in its middle. It after calls `balance(std::vector<std::pair<T, list, Iterator>>, list, Iterator begin, Iterator end)` on the halves recursively, till there are no more nodes to insert. The `const` operator `[]` return a `std::runtime_exception` if a non valid key is accessed. There are also some test methods which are defined and declared only when loaded in the unit test that find the height of the tree and weather it is balanced or not. The unit test makes use of the `catch2` unit testing framework, which is a collection of macros stored in the `catch.hpp` file and speeds up the writing of the test code.

Performance measures

In order to evaluate the performances of our implementation of the binary tree, we ran some tests. The test involves four different types of tree: a "linked-list" tree, created inserting elements orderly (with respect to the key), a random tree, created inserting elements with keys were in random order, a balanced tree, created copying and balancing the random tree, and finally the `std::map`. Since the performances of the linked-list-tree were much worse than the one of the other trees, we decided to ran the test for different sizes for these two different cases, in this way we can run significant tests for the faster trees. The benchmark consists (for each tree) in accessing all the elements and measuring the average access time: In the case of the linked-list-tree we expect to see a time that is $O(n)$ with respect to the tree size, for the balanced and the `std::map` we expect to see $O(\log_2 n)$. The results are shown in microseconds, and code is compiled with the level `O3` of optimisation.

As said before the linked-list-tree is much slower, about 100 times slower for a tree size that is 10 times smaller. From the plots it seems that the previous considerations about the time complexity hold. We can see that even though in this case balancing the tree improves dramatically the performances (probably by a factor greater than 100), also the random tree is quite good and in this case the speedup is less than 2.